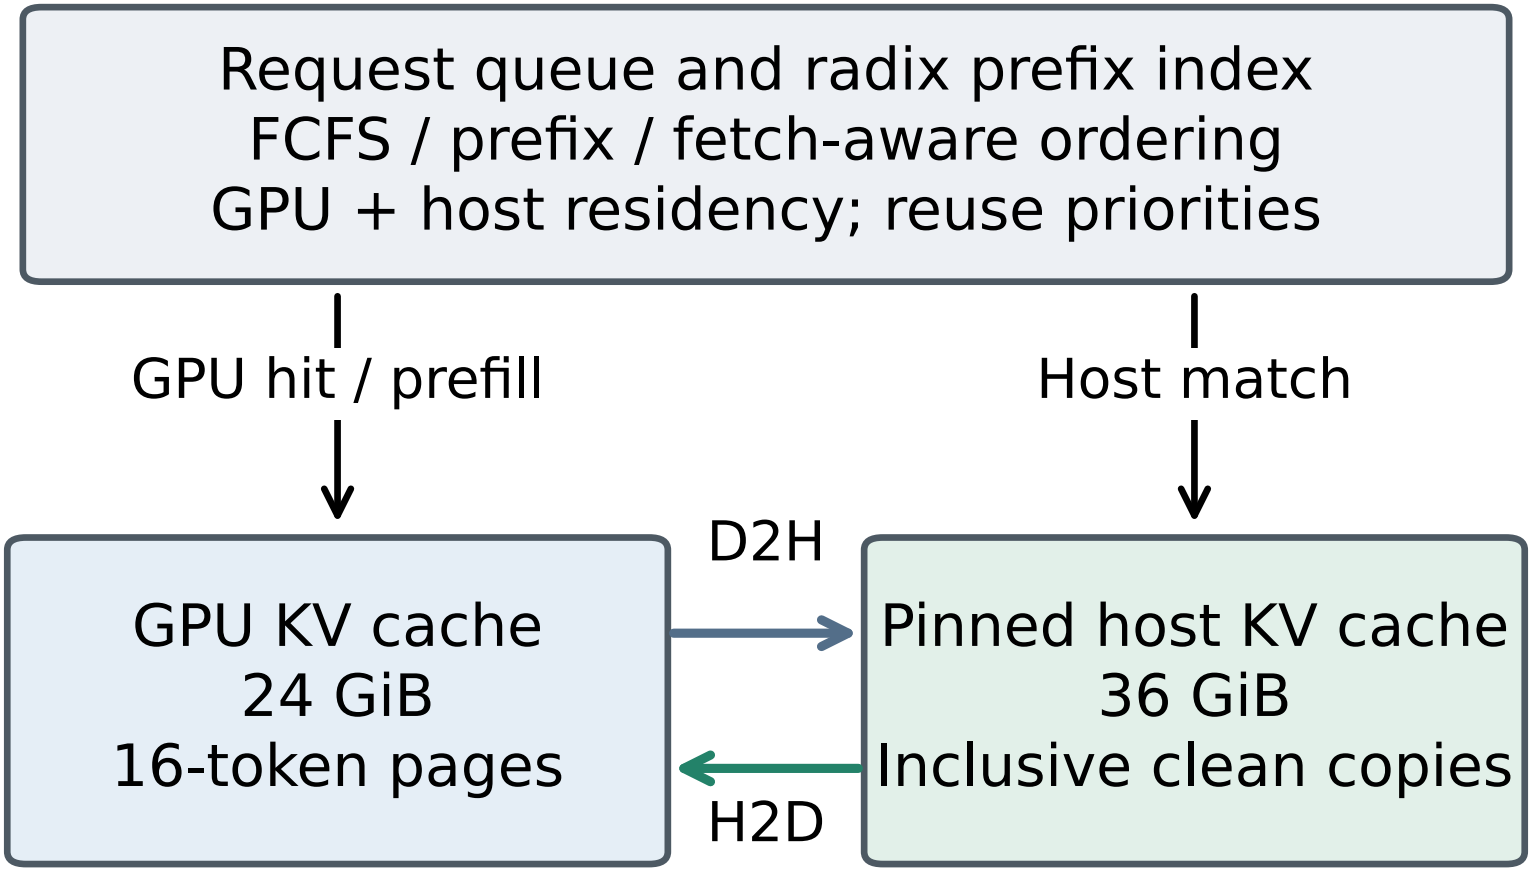
Request queue and radix prefix index
FCFS / prefix / fetch-aware ordering
GPU + host residency; reuse priorities
GPU hit / prefill
Host match
GPU KV cache
24 GiB
16-token pages
D2H
H2D
Pinned host KV cache
36 GiB
Inclusive clean copies
Dirty eviction: offload | Reuse: restore / prefetch