INTRODUCTION

CacheWise replication and KV cache recovery

Replicating CacheWise and evaluating host-DRAM recovery

CENTRAL QUESTION

How can coding agents reuse previous computation when GPU memory is limited?

What do the CacheWise traces contain?

CATraces: real coding sessions

Claude Code interaction logs
Model requests + tool calls
Timing + prompt / output lengths

Released content is redacted; replay
preserves lengths and prefix growth.

Reuse across turns

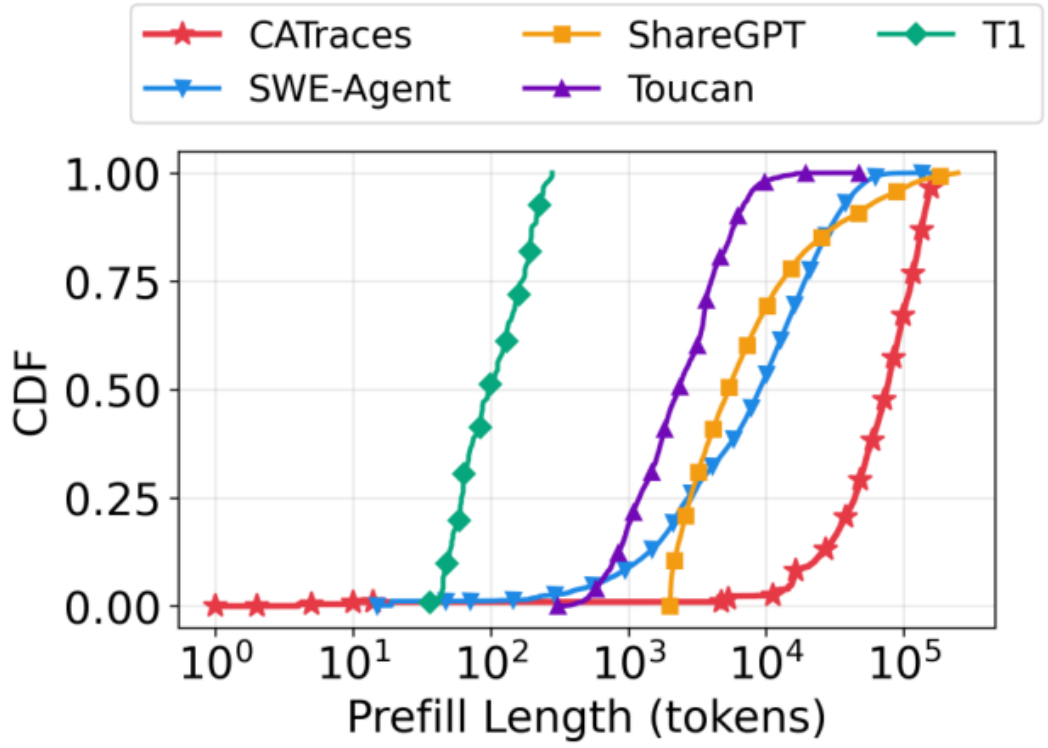
Growing, shared prompt prefixes
Long inputs with short outputs
Tool waits between requests

Requests arrive in sessions,
with repeated opportunities for reuse.

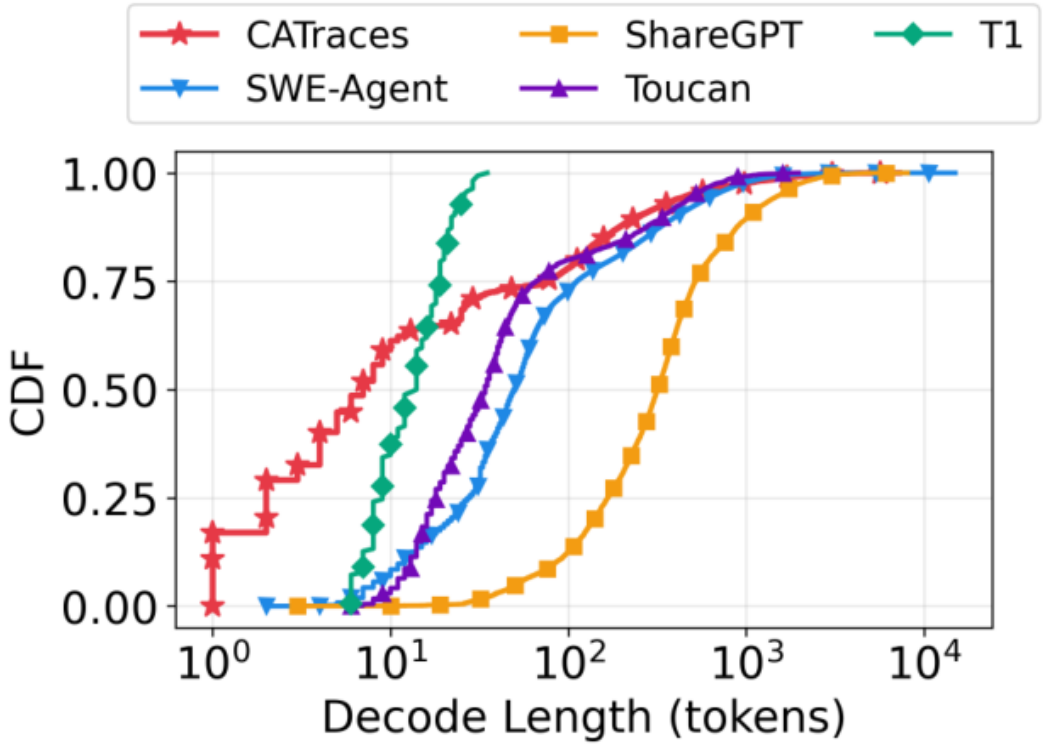
Why prefix-aware scheduling can help

Reuse a resident prefix → less new prefill → fewer additional KV blocks.
A smaller active working set can reduce eviction pressure.

CATraces: long inputs, short outputs



(a) Prefill length distributions for different datasets



(b) Decode length distributions for different datasets

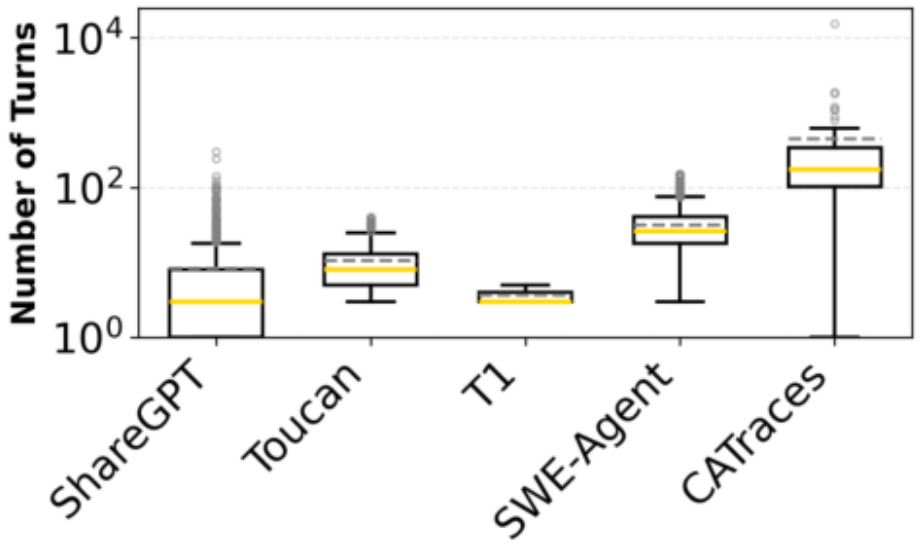
CATraces (red) vs. ShareGPT chat (orange): longer prefill, shorter decode.

CDF = fraction of requests at or below the token length. Published results, not our replay.

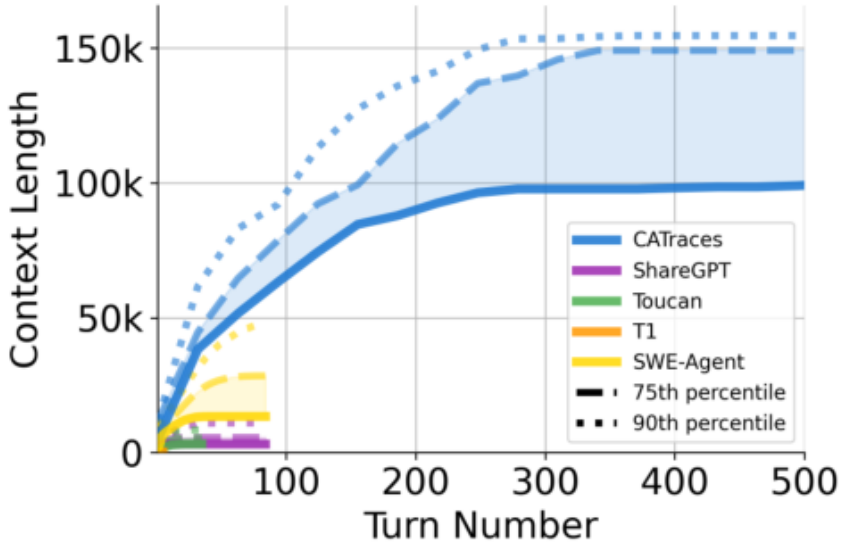
Cropped from Tiwari et al., CacheWise v1, Fig. 2 (PDF p. 4); CC BY 4.0

CATraces differ from traditional chat

More turns per session (Fig. 3)



Growing context across turns (Fig. 6)



One task drives many model/tool turns; growing KV prefixes compete for memory.
Session completion matters alongside individual request latency.

Cropped from Tiwari et al., CacheWise v1, Figs. 3 and 6 (PDF p. 4); CC BY 4.0

CacheWise predicts when a session returns

Schedule requests

Additional KV blocks = required − resident

Serve the smallest additional-block cost first.

Uses current GPU prefix overlap.

Predict duration → order evictions

Tool name + arguments → **TF–IDF → trained KMeans Historical duration samples**

Condition on the tool still running:

$$r = E[D - e \mid D > e]$$

Evict unreferenced KV with largest r; refresh predictions as time passes.

D = total duration; e = elapsed time. Replication: tool names only (arguments redacted).

Algorithm adapted from CacheWise v1 §§5.1–5.3; original-paper argument-aware predictor

CacheWise combines scheduling and eviction

Scheduler (rows)	Eviction policy (columns)	
	LRU	**Predictive**
FCFS	**Baseline** **FCFS + LRU**	**Eviction only** **FCFS + predictive**
Prefix-aware	**Scheduling only** **Prefix + LRU**	**CacheWise (both)** **Prefix + predictive**

Bottom-right: adapted CacheWise uses both components.

Testbed

Qwen2.5-Coder-14B-Instruct

1 × RTX PRO 6000 Blackwell Max-Q GPU

BF16 · tensor parallelism 1

KV: 287,696 tokens (17,981 × 16)

Prefill chunk: 512 tokens

Max sequence length: 131,072

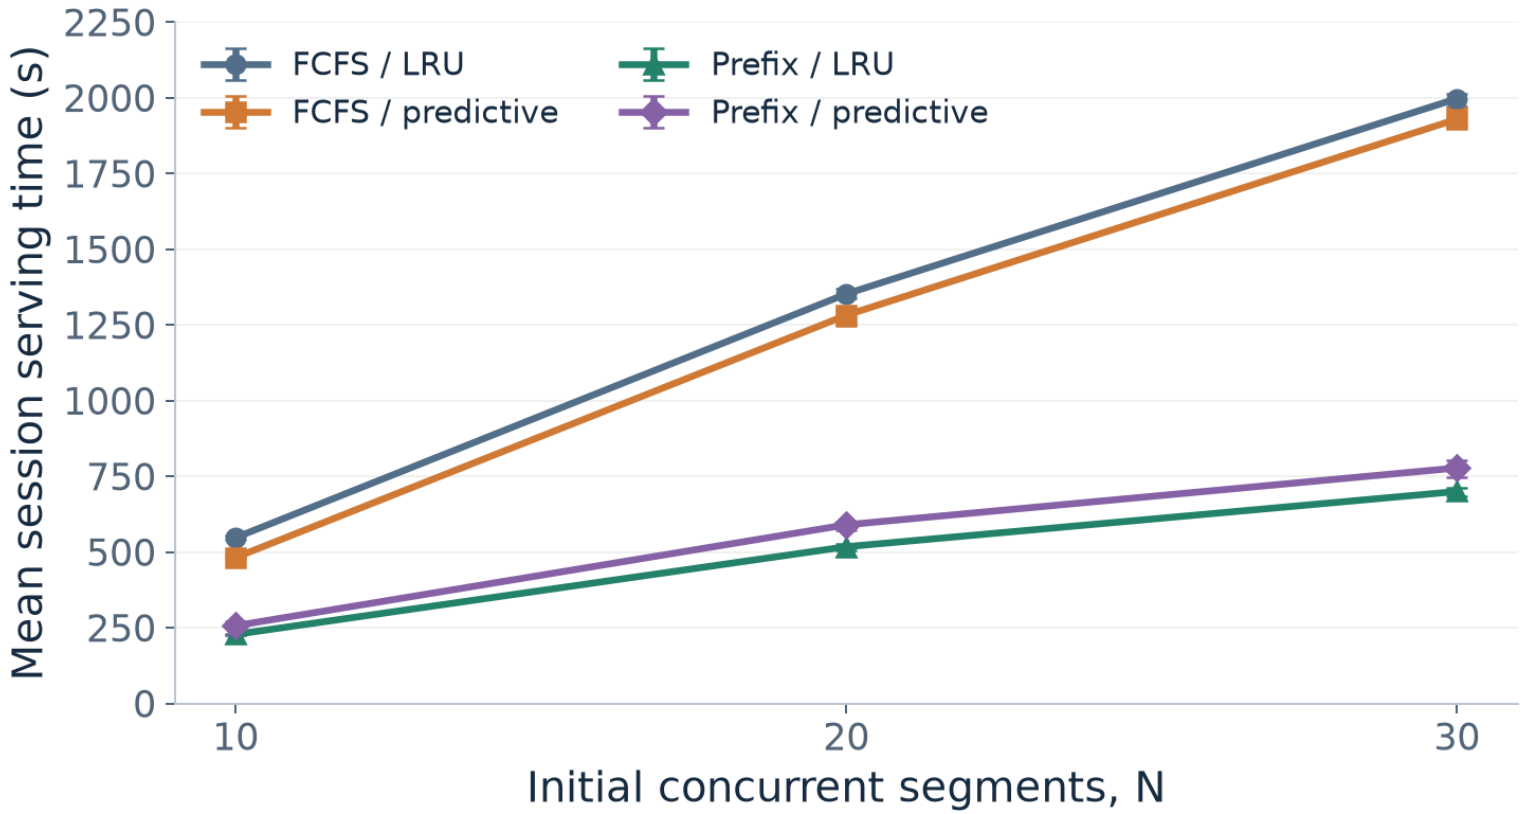
Max scheduled sequences: 32

Coverage: 5 settings × 4 policies × 3 repeats = 60 runs · 6,384 requests

Metric: mean summed model-request latency per session, excluding tool waits.

Evidence: four-policy replication; report Tables 1–3 and §3

Prefix / LRU has the lowest mean



5 / 5

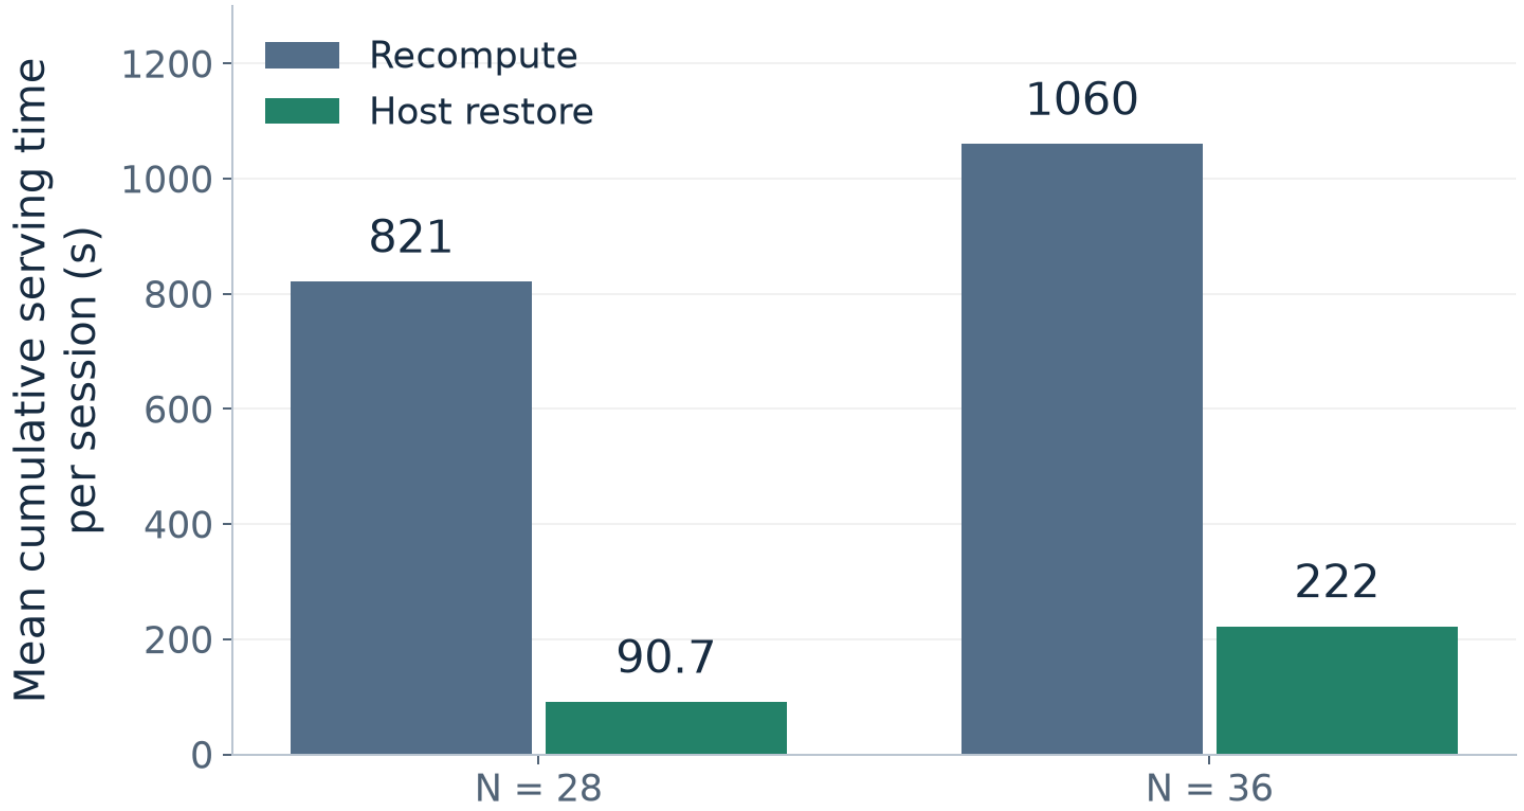
settings favor prefix / LRU

58–65% lower mean vs. FCFS / LRU in the load sweep

Three repeats per cell. Ranges are observed minima/maxima; lower is better.

Host restoration reduces session time

Recovery tradeoff: save + restore + management < recomputation



9.1× / 4.8×

reported time ratios at N = 28 / 36

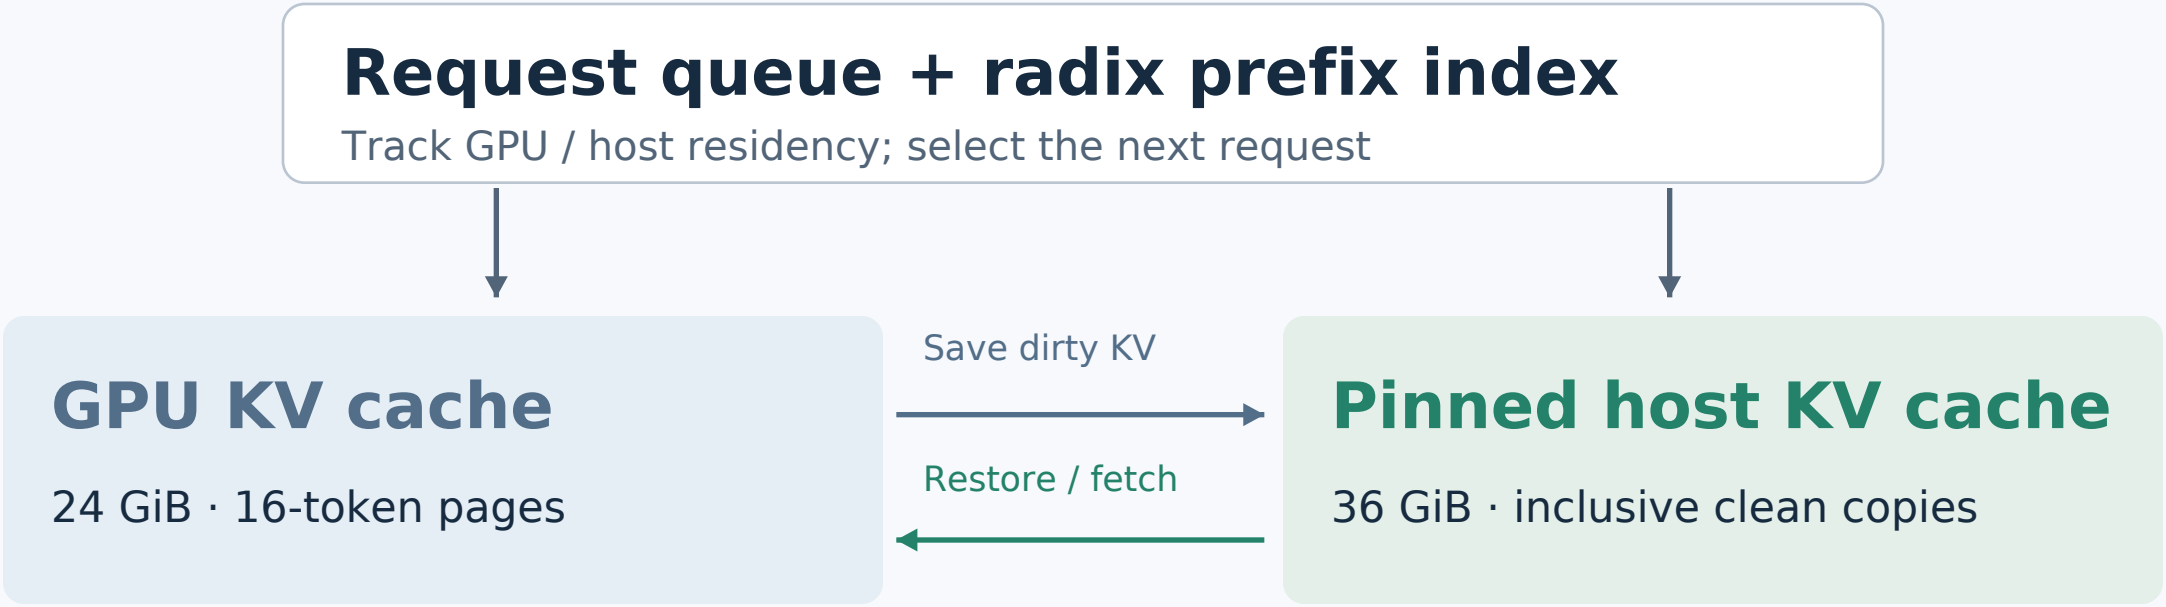
+36 GiB host KV
24 GiB GPU KV fixed

One observation per bar

Separate Mini-SGLang study. Added-capacity comparison; lower is better.

Evidence: host-dram-reported-results.json; report Figure 3a and §6

Preserve evicted KV in host DRAM



On demand

Restore before prefill.

Prefetch

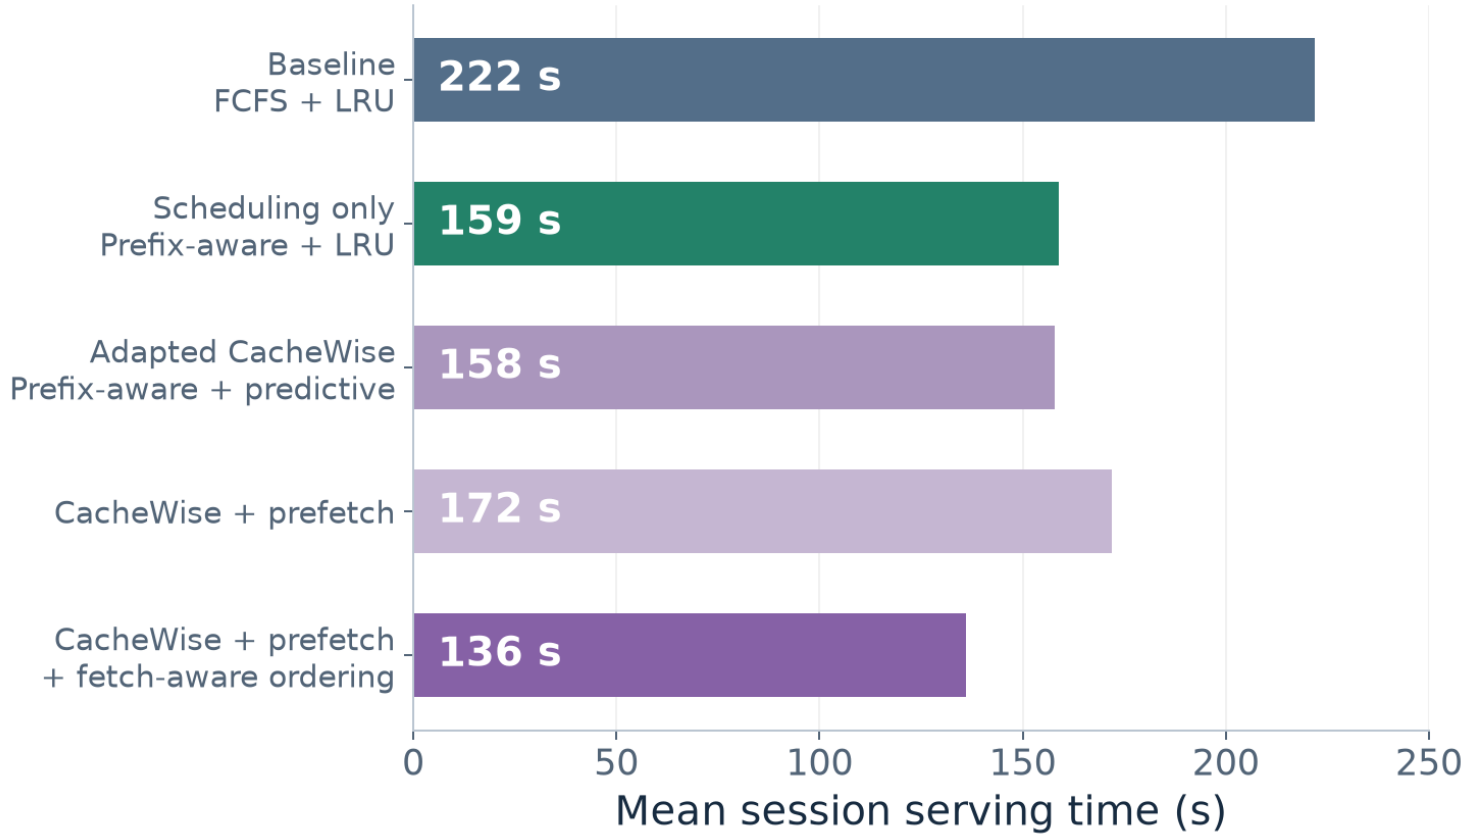
Overlap predicted reuse.

Fallback

Recompute on a miss.

Fetch-aware ordering discounts host hits

Fetch-aware priority = new prefill tokens + 0.05 × additional host tokens



What changes?

Prefix-aware

Order by GPU reuse;
less new prefill first.

Prefetch

Restore host KV
before session return.

Fetch-aware

Also discount host hits
when ordering requests.

N = 36, host restoration in every row. First reported runs; sparse repetitions.

Evidence: Mini-SGLangs 3f9aee7; report Figure 3b and Eq. (3)

Keep the conclusions within the evidence

Replication

Synthetic replay and selected segments.

One retained preemption; extra prefill work unmeasured.

Host results

Separate engine and replay.

Added host capacity; most configurations run once.

SSD extension

Not implemented or measured.

Needs full I/O, staging, metadata and recovery.

Interpret both studies separately; retain the declared protocol deviation.

Three takeaways

01

Scheduling and eviction interact.

Prefix / LRU has the lowest mean in all five replication settings.

02

Host copies can avoid costly recovery.

Large reported gains include the added 36 GiB backing tier.

03

Order requests by recovery cost.

Fetch-aware results are promising; SSD evaluation remains future work.